

# Designing a Flash Translation Layer

19 February, 2016  
Niv Dayan

# Introduction

- Niv Dayan (post-doc)
  - Data Systems Lab
  - IACS
- Today's talk:
  - Designing a Flash Translation Layer
- Hopeful learning outcomes
  - What flash looks like internally
  - See a cool application of LSM-trees



# Introduction

- Structure
  - Background (how flash works)
  - Scalability Problem (for terabyte flash devices)
  - Solution (LSM-tree variant)
- Questions/discussions are welcome



# Background

# Background

- Flash devices increasingly used for many applications
- Advantages:
  - Good read/write performance
  - Low power consumption
- Internally they are very complex




# Background




- Ideosyncracies
  - Reads & writes at page granularity
  - Sequential writes within a block
  - Block-erase before update
  - Limited erases per block

# Background




- Read/write **asymmetry**
  - Reads: tens of microseconds
  - Writes: few milliseconds




# Background

- A flash translation layer (FTL) hides these constraints
- Exposes a simple block interface to Application




# Background

- FTL writes pages sequentially within a block
- Recall erase-before-write rule
- Reduce cost through out-of-place updates




# Background

- Eventually, invalid pages accumulate
- Perform garbage-collection to free space




# Background

- Eventually, invalid pages accumulate
- Perform garbage-collection to free space




# Background

- Logical pages move around
- The FTL must maintain a mapping table from logical to physical addresses
- Device has RAM for storing metadata



# Background


- **Example**
  - Application reads logical page 300



# Background

- **Problem**
  - The mapping table is large
  - For a 2 TB flash device, 2 GB mapping table
  - Not enough integrated RAM


- Solution?
  - Store it in flash



Integrated  
RAM

# Background

- Mapping table is stored in flash
- E.g. Application reads page 300
- We pay in performance



# Background

- To reduce the overhead of mapping reads, we use a cache

Integrated RAM

Logical mapping range	flash page
256 - 383	47

Cache

Logical address	Physical address
300	X
...	...


Flash

Mapping pages	Data pages
...	...
...	300
47	...

# Background

- E.g. suppose application updates page 300
- We must update mapping page 47
- This can harm performance.
- Mark entry as “dirty”. Update lazily later.

Integrated RAM




Flash

# Background

- E.g. suppose application updates page 300
- We must update mapping page 47
- This can harm performance.
- Mark entry as “dirty”. Update lazily later.

Integrated RAM




Flash


# Background

- E.g. suppose application updates page 300
- We must update mapping page 47
- This can harm performance.
- Mark entry as “dirty”. Update lazily later.

Integrated RAM



Flash



# Background

- The FTL also maintains a bitmap
- Keeps track of which pages are invalid
- Needed for garbage-collection

Integrated RAM

Mapping  
directory

Cache

Page  
Validity  
Bitmap  
(PVB)

# Background

- PVB enables garbage-collection victim-selection
- Which pages to migrate

Integrated RAM

PVB section  
for block X

...	0	0	0	0	0	1	0	0	...
-----	---	---	---	---	---	---	---	---	-----


Flash

Flash block X

10
21
1
3
24
18
4
19


# Background

- PVB enables garbage-collection victim-selection
- Which pages to migrate



# Background

- PVB enables garbage-collection victim-selection
- Which pages to migrate



# Background

- PVB enables garbage-collection victim-selection
- Which pages to migrate

Integrated RAM

PVB section  
for block X

...	0	0	1	0	0	1	0	0	...
-----	---	---	---	---	---	---	---	---	-----

Flash

Flash block X


Garbage-collected

10
21
1
3
24
18
4
19


# Background

- PVB enables garbage-collection victim-selection
- Which pages to migrate

Integrated RAM




Flash



# Background

- PVB enables garbage-collection victim-selection
  - Which pages to migrate



# Background

- Summary
  - Logical pages are updated out-of-place
  - Garbage-collection takes place to free space
  - A translation table keeps track of location of data
  - The translation table is stored in flash
  - PVB is used to identify live pages in GC
  - PVB is stored in integrated RAM

# Background

- Integrated RAM is volatile
- When power fails, we lose all RAM-resident metadata
- How can we recover PVB
- Answer: scan all mapping pages



Integrated RAM


Mapping  
directory

Cache


Page  
Validity  
Bitmap  
(PVB)

Flash

Mapping  
pages




Data  
pages



# Problem


# Problem

- Size of PVB is proportional to capacity
- As flash devices scale to terabytes, size of PVB is becoming impractical



# Problem

- Integrated RAM is volatile
- When power fails, we lose PVB.
- Recovering PVB is too long



# Problem

- **Observations:**
  - PVB consumes 95% of integrated RAM
  - Recovering PVB takes 51% of recovery time
- **Simple solution**
  - Store PVB in flash
  - Problem: bad for performance



# Problem

- Why?
  - Every time a logical page is updated, we must also update PVB.
  - This doubles the number of writes

Integrated RAM

PVB section  
for block X

...	0	0	0	0	0	1	0	0	...
-----	---	---	---	---	---	---	---	---	-----

Flash

# Problem

- Recall that writes are 10-100 times more expensive than reads on flash
- Writes also wear out the device
- Thus, storing PVB naively in flash:
  - halves throughput
  - halves device lifetime

# Problem

- **Definition:** Write-amplification is the number of physical writes taking place internally for each logical write.
- **Goal:**
  - Store PVB in flash
  - Keep write-amplification low
  - Fast lookup time

# Solution

# Solution

- Insight:
  - PVB is updated once per application write
  - PVB is accessed once per garbage-collection operation
  - A garbage-collection operation happens once for every  $\approx 100$  application writes


# Solution

- Since PVB is mostly updated, we propose a write-optimized data structure
- We use an LSM-tree with some modifications

# Solution


- Key ideas: buffer updates about page validity in integrated RAM
- When the buffer fills, flush to flash
- Reorganize this data in flash periodically to keep access time fast

# Solution




# Solution

- When a page is invalidated, create entry:
  - Block ID that invalidated page is in
  - Bitmap with B bits set to 0
  - Mark invalidated page as 1




- Entry is inserted into the buffer

# Solution




# Solution




# Solution

- When a block is erased, create an entry
  - Block ID of erased block
  - All bits set to 0
  - Erase flag set to 1




- Entry is inserted into the buffer


# Solution




# Solution




# Solution




# Solution




# Solution




# Solution




# Solution




# Solution




# Solution




# Solution




# Solution



# Solution




# Solution




# Solution

- Merge bitmaps using bitwise “or” operation




# Solution


- However, if erase flag of more recent entry is 1, ignore less recent entry




# Solution




# Solution




# Solution




# Solution




# Solution




# Solution




# Solution



# Solution




# Solution




# Solution

- Garbage-collect block 4




# Solution

- Garbage-collect block 4




# Solution

- Garbage-collect block 4




# Solution

- Garbage-collect block 4




# Solution

- Garbage-collect block 4




# Solution

- Garbage-collect block 4




# Solution

- Garbage-collect block 4




# Solution

- Garbage-collect block 4



# Solution

- Garbage-collect block 4



# Solution

- **Lookup:** In general, search all substructures from smallest to largest, or until we find entry with erase flag set to 1
- Merge bitmaps
- We now have an image of which pages in the block are invalid
- Garbage-collection can take place

# Analysis

# Analysis

- K: number of flash blocks
- V: number of entries in the buffer
- T: LSM-tree's size ratio
- The number of levels is:  $O(\log_T(\frac{K}{V}))$

# Analysis

- Each entry is copied  $O(T)$  times within a level
- Each entry is copied  $O(\log_T(\frac{K}{V}))$  times across levels
- Copying an entry takes  $O(\frac{1}{V})$  flash write
- Write-amplification:  $O(\frac{T}{V} \log_T(\frac{K}{V}))$


# Analysis

- This is sub-constant  $O\left(\frac{T}{V} \log_T\left(\frac{K}{V}\right)\right)$
- For typical values  $\approx 0.02$
- During garbage-collection  $O\left(\log\left(\frac{K}{V}\right)\right)$  reads
- But reads are cheap in flash
- And garbage-collection happens infrequently

# Evaluation


# Analysis

- Tuning the LSM-tree's size ratio




# Analysis

- Write-amplification is low and scales




# Analysis

- RAM-requirement is 95% lower



# Analysis

- Recovery time is 51% lower



# Other Concerns

# Other Concerns

- **Power Failure**
  - We lose the contents of the buffer
  - We lose dirty mapping entries in the cache

Integrated RAM

Mapping  
directory

Cache

LSM-Buffer

# Other Concerns

- **Picking Garbage-Collection Victims**
  - Greedy approach: least number of live pages
  - Non-optimal

Candidate 1

10
21
1
3
24
18
4
19

Candidate 2

53
31
52
16
74
85
43
32

# Other Concerns

- **Reason:** suppose blocks in block 2 are extremely frequently updated.
- Most would soon be invalidated anyways.

Candidate 1


10
21
1
3
24
18
4
19

Candidate 2

53
31
52
16
74
85
43
32

# Other Concerns


- Write-amplification depends on V:  $O(\frac{T}{V} \log_T(\frac{K}{V}))$
- V is the number of entries fitting into buffer
- Entry size depends on block size



- As devices scale, block size increases, so V decreases

# Other Concerns

- Solution: entry-partitioning



# Conclusion

# Conclusion

- We can store page validity metadata in flash while keeping write-amplification low
- Reduces RAM requirement by 95%
- Reduces recovery time by 51%
- Write-amplification increases by 0.02%

# Conclusion

- LSM-trees can be applied beyond their usual application for key-value storage
- Thanks. Q&A.